

• MobaXterm Personal Edition v22.0 •
(X server, SSH client and network tools)

- Your computer drives are accessible through the `/drives` path
- Your DISPLAY is set to `192.xxx.x.xxx:x.x`
- When using `SSH`, your remote DISPLAY is automatically forwarded
- Each command status is specified by a special symbol (`v` or `x`)

• Important:

This is MobaXterm Personal Edition. The Professional edition allows you to customize MobaXterm for your company: you can add your own logo, your parameters, your welcome message and generate either an MSI installation package or a portable executable. We can also modify MobaXterm or develop the plugins you need. For more information: <https://mobaxterm.mobatek.net/download.html>

```
24/04/2022 19:08.20 /home/mobaxterm ssh 192.xxx.x.xxx:x.x -l everis
```

```
everis@192.xxx.x.xxx:x.x 's password:
```

```
X11 forwarding request failed on channel 0
```

```
Last login: Sun Apr 24 20:07:41 2022
```

```
[everis@bigdata-srv ~]$ hbase shell
```

```
2022-04-24 20:33:56,674 INFO [main] Configuration.deprecation: hadoop.native.lib is deprecated. Instead, use io.native.lib.available
```

```
HBase Shell; enter 'help<RETURN>' for list of supported commands.
```

```
Type "exit<RETURN>" to leave the HBase Shell
```

```
Version 1.2.0-cdh5.16.2, rUnknown, Mon Jun 3 03:50:03 PDT 2019
```

```
hbase(main):001:0> status
```

```
ERROR: org.apache.hadoop.hbase.ipc.ServerNotRunningYetException: Server is not running yet
```

```
at org.apache.hadoop.hbase.master.HMaster.checkServiceStarted(HMaster.java:2404)
```

```
at org.apache.hadoop.hbase.master.MasterRpcServices.isMasterRunning(MasterRpcServices.java:941)
```

```
at org.apache.hadoop.hbase.protobuf.generated.MasterProtos$MasterService$2.callBlockingMethod(MasterProtos.java:58686)
```

```
at org.apache.hadoop.hbase.ipc.RpcServer.call(RpcServer.java:2191)
```

```
at org.apache.hadoop.hbase.ipc.CallRunner.run(CallRunner.java:112)
```

```
at org.apache.hadoop.hbase.ipc.RpcExecutor$Handler.run(RpcExecutor.java:183)
```

```
at org.apache.hadoop.hbase.ipc.RpcExecutor$Handler.run(RpcExecutor.java:163)
```

```
Show cluster status. Can be 'summary', 'simple', 'detailed', or 'replication'. The
```

```
default is 'summary'. Examples:
```

```
hbase> status
```

```
hbase> status 'simple'
```

```
hbase> status 'summary'
```

```
hbase> status 'detailed'
```

```
hbase> status 'replication'
```

```
hbase> status 'replication', 'source'
```

```
hbase> status 'replication', 'sink'
```

```
hbase(main):002:0> status simple
NameError: undefined local variable or method `simple' for #<Object:0x44e4cb76>
```

```
hbase(main):003:0> status 'simple'
```

```
ERROR: org.apache.hadoop.hbase.ipc.ServerNotRunningYetException: Server is not running yet
    at org.apache.hadoop.hbase.master.HMaster.checkServiceStarted(HMaster.java:2404)
    at
org.apache.hadoop.hbase.master.MasterRpcServices.isMasterRunning(MasterRpcServices.java:94
1)
    at
org.apache.hadoop.hbase.protobuf.generated.MasterProtos$MasterService$2.callBlockingMethod
(MasterProtos.java:58686)
    at org.apache.hadoop.hbase.ipc.RpcServer.call(RpcServer.java:2191)
    at org.apache.hadoop.hbase.ipc.CallRunner.run(CallRunner.java:112)
    at org.apache.hadoop.hbase.ipc.RpcExecutor$Handler.run(RpcExecutor.java:183)
    at org.apache.hadoop.hbase.ipc.RpcExecutor$Handler.run(RpcExecutor.java:163)
```

Show cluster status. Can be 'summary', 'simple', 'detailed', or 'replication'. The default is 'summary'. Examples:

```
hbase> status
hbase> status 'simple'
hbase> status 'summary'
hbase> status 'detailed'
hbase> status 'replication'
hbase> status 'replication', 'source'
hbase> status 'replication', 'sink'
```

```
hbase(main):004:0> status summary
NameError: undefined local variable or method `summary' for #<Object:0x44e4cb76>
```

```
hbase(main):005:0> status 'summary'
```

```
ERROR: org.apache.hadoop.hbase.ipc.ServerNotRunningYetException: Server is not running yet
    at org.apache.hadoop.hbase.master.HMaster.checkServiceStarted(HMaster.java:2404)
    at
org.apache.hadoop.hbase.master.MasterRpcServices.isMasterRunning(MasterRpcServices.java:94
1)
    at
org.apache.hadoop.hbase.protobuf.generated.MasterProtos$MasterService$2.callBlockingMethod
(MasterProtos.java:58686)
    at org.apache.hadoop.hbase.ipc.RpcServer.call(RpcServer.java:2191)
    at org.apache.hadoop.hbase.ipc.CallRunner.run(CallRunner.java:112)
    at org.apache.hadoop.hbase.ipc.RpcExecutor$Handler.run(RpcExecutor.java:183)
    at org.apache.hadoop.hbase.ipc.RpcExecutor$Handler.run(RpcExecutor.java:163)
```

Show cluster status. Can be 'summary', 'simple', 'detailed', or 'replication'. The default is 'summary'. Examples:

```
hbase> status
hbase> status 'simple'
hbase> status 'summary'
hbase> status 'detailed'
hbase> status 'replication'
hbase> status 'replication', 'source'
```

```
hbase> status 'replication', 'sink'
```

```
hbase(main):006:0> version  
1.2.0-cdh5.16.2, rUnknown, Mon Jun  3 03:50:03 PDT 2019
```

```
hbase(main):007:0> status 'detailed'
```

```
ERROR: org.apache.hadoop.hbase.ipc.ServerNotRunningYetException: Server is not running yet  
    at org.apache.hadoop.hbase.master.HMaster.checkServiceStarted(HMaster.java:2404)  
    at org.apache.hadoop.hbase.master.MasterRpcServices.isMasterRunning(MasterRpcServices.java:941)  
    at org.apache.hadoop.hbase.protobuf.generated.MasterProtos$MasterService$2.callBlockingMethod(MasterProtos.java:58686)  
    at org.apache.hadoop.hbase.ipc.RpcServer.call(RpcServer.java:2191)  
    at org.apache.hadoop.hbase.ipc.CallRunner.run(CallRunner.java:112)  
    at org.apache.hadoop.hbase.ipc.RpcExecutor$Handler.run(RpcExecutor.java:183)  
    at org.apache.hadoop.hbase.ipc.RpcExecutor$Handler.run(RpcExecutor.java:163)
```

Show cluster status. Can be 'summary', 'simple', 'detailed', or 'replication'. The default is 'summary'. Examples:

```
hbase> status  
hbase> status 'simple'  
hbase> status 'summary'  
hbase> status 'detailed'  
hbase> status 'replication'  
hbase> status 'replication', 'source'  
hbase> status 'replication', 'sink'
```

```
hbase(main):008:0> status 'replication'
```

```
ERROR: org.apache.hadoop.hbase.ipc.ServerNotRunningYetException: Server is not running yet  
    at org.apache.hadoop.hbase.master.HMaster.checkServiceStarted(HMaster.java:2404)  
    at org.apache.hadoop.hbase.master.MasterRpcServices.isMasterRunning(MasterRpcServices.java:941)  
    at org.apache.hadoop.hbase.protobuf.generated.MasterProtos$MasterService$2.callBlockingMethod(MasterProtos.java:58686)  
    at org.apache.hadoop.hbase.ipc.RpcServer.call(RpcServer.java:2191)  
    at org.apache.hadoop.hbase.ipc.CallRunner.run(CallRunner.java:112)  
    at org.apache.hadoop.hbase.ipc.RpcExecutor$Handler.run(RpcExecutor.java:183)  
    at org.apache.hadoop.hbase.ipc.RpcExecutor$Handler.run(RpcExecutor.java:163)
```

Show cluster status. Can be 'summary', 'simple', 'detailed', or 'replication'. The default is 'summary'. Examples:

```
hbase> status
hbase> status 'simple'
hbase> status 'summary'
hbase> status 'detailed'
hbase> status 'replication'
hbase> status 'replication', 'source'
hbase> status 'replication', 'sink'
```

```
hbase(main):009:0> table_help
Help for table-reference commands.
```

You can either create a table via 'create' and then manipulate the table via commands like 'put', 'get', etc.
See the standard help information for how to use each of these commands.

However, as of 0.96, you can also get a reference to a table, on which you can invoke commands.
For instance, you can get create a table and keep around a reference to it via:

```
hbase> t = create 't', 'cf'
```

Or, if you have already created the table, you can get a reference to it:

```
hbase> t = get_table 't'
```

You can do things like call 'put' on the table:

```
hbase> t.put 'r', 'cf:q', 'v'
```

which puts a row 'r' with column family 'cf', qualifier 'q' and value 'v' into table t.

To read the data out, you can scan the table:

```
hbase> t.scan
```

which will read all the rows in table 't'.

Essentially, any command that takes a table name can also be done via table reference.

Other commands include things like: get, delete, deleteall, get_all_columns, get_counter, count, incr. These functions, along with the standard JRuby object methods are also available via tab completion.

For more information on how to use each of these commands, you can also just type:

```
hbase> t.help 'scan'
```

which will output more information on how to use that command.

You can also do general admin actions directly on a table; things like enable, disable, flush and drop just by typing:

```
hbase> t.enable
```

```
hbase> t.flush
hbase> t.disable
hbase> t.drop
```

Note that after dropping a table, your reference to it becomes useless and further usage is undefined (and not recommended).

```
hbase(main):010:0> list
TABLE
```

```
ERROR: org.apache.hadoop.hbase.ipc.ServerNotRunningYetException: Server is not running yet
    at org.apache.hadoop.hbase.master.HMaster.checkServiceStarted(HMaster.java:2404)
    at org.apache.hadoop.hbase.master.MasterRpcServices.isMasterRunning(MasterRpcServices.java:941)
    at org.apache.hadoop.hbase.protobuf.generated.MasterProtos$MasterService$2.callBlockingMethod(MasterProtos.java:58686)
    at org.apache.hadoop.hbase.ipc.RpcServer.call(RpcServer.java:2191)
    at org.apache.hadoop.hbase.ipc.CallRunner.run(CallRunner.java:112)
    at org.apache.hadoop.hbase.ipc.RpcExecutor$Handler.run(RpcExecutor.java:183)
    at org.apache.hadoop.hbase.ipc.RpcExecutor$Handler.run(RpcExecutor.java:163)
```

List all tables in hbase. Optional regular expression parameter could be used to filter the output. Examples:

```
hbase> list
hbase> list 'abc.*'
hbase> list 'ns:abc.*'
hbase> list 'ns:.*'
```

```
hbase(main):011:0> ls
NameError: undefined local variable or method `ls' for #<Object:0x44e4cb76>
```

```
hbase(main):012:0> cd ..
hbase(main):013:0* ls
NameError: undefined local variable or method `cd' for #<Object:0x44e4cb76>
```

```
hbase(main):014:0> logout
NameError: undefined local variable or method `logout' for #<Object:0x44e4cb76>
```

```
hbase(main):015:0> exit
[everis@bigdata-srv ~]$ ls
000000_0          employee.txt          latitude            sq_fit
avengers.csv      file_teste.txt       longitude          state
base_localidade.csv  firewall.rules       metastore_db       street
baths            FL_insurance_sample.csv  pokemon.sql        tmp
beds             hive-site.xml        price              type
city            import_sqoop.sh      sale_date          zip
derby.log        install_sqoop.sh     script_apoio
[everis@bigdata-srv ~]$ cd script_apoio/
[everis@bigdata-srv script_apoio]$ ls
cassandra_start_all_service.sh  restart_all_service.sh
```

```
cassandra_stop_all_service.sh  start_all_service.sh
output.txt                     stop_all_service.sh
restart_all_service_hive.sh
```

```
[everis@bigdata-srv script_apoio]$ sudo -u hdfs hadoop fsck / | egrep -v '^\.+$'
| grep -v eplica
```

```
[sudo] password for everis:
```

```
DEPRECATED: Use of this script to execute hdfs command is deprecated.
Instead use the hdfs command for it.
```

```
Connecting to namenode via http://bigdata-srv:50070/fsck?ugi=hdfs&path=%2F
FSCK started by hdfs (auth:SIMPLE) from /192.xxx.x.xxx:x.x for path / at Sun Apr 24
21:08:49 PMDT 2022
/tmp/wordcount00002.out/part-r-00000: CORRUPT blockpool BP-295589643-10.0.2.15-1
610551694034 block blk_1073742421
```

```
/tmp/wordcount00002.out/part-r-00000: MISSING 1 blocks of total size 529 B...
/tmp/wordcount012.out/part-r-00000: CORRUPT blockpool BP-295589643-10.0.2.15-161
0551694034 block blk_1073742361
```

```
/tmp/wordcount012.out/part-r-00000: MISSING 1 blocks of total size 529 B..
```

```
/user/history/done_intermediate/hdfs/job_1611095949321_0001.summary: CORRUPT blo
ckpool BP-295589643-10.0.2.15-1610551694034 block blk_1073742422
```

```
/user/history/done_intermediate/hdfs/job_1611095949321_0001.summary: MISSING 1 b
locks of total size 347 B..
```

```
/user/history/done_intermediate/root/job_1611089809353_0001-1611091594201-root-w
ord+count-1611091609627-1-1-SUCCEEDED-root.root-1611091600450.jhist: CORRUPT blo
ckpool BP-295589643-10.0.2.15-1610551694034 block blk_1073742363
```

```
/user/history/done_intermediate/root/job_1611089809353_0001-1611091594201-root-w
ord+count-1611091609627-1-1-SUCCEEDED-root.root-1611091600450.jhist: MISSING 1 b
locks of total size 33666 B..
```

```
/user/history/done_intermediate/root/job_1611089809353_0001.summary: CORRUPT blo
ckpool BP-295589643-10.0.2.15-1610551694034 block blk_1073742362
```

```
/user/history/done_intermediate/root/job_1611089809353_0001.summary: MISSING 1 b
locks of total size 347 B..
```

```
/user/history/done_intermediate/root/job_1611089809353_0001_conf.xml: CORRUPT bl
ockpool BP-295589643-10.0.2.15-1610551694034 block blk_1073742364
```

```
/user/history/done_intermediate/root/job_1611089809353_0001_conf.xml: MISSING 1
blocks of total size 141681 B..
```

```
/var/log/hadoop-yarn/apps/hdfs/logs/application_1611095949321_0001/bigdata-srv_3
4788: CORRUPT blockpool BP-295589643-10.0.2.15-1610551694034 block blk_107374242
5
```

```
/var/log/hadoop-yarn/apps/hdfs/logs/application_1611095949321_0001/bigdata-srv_3
4788: MISSING 1 blocks of total size 43862 B..
```

```
/var/log/hadoop-yarn/apps/root/logs/application_1611089809353_0001/bigdata-srv_4
2232: CORRUPT blockpool BP-295589643-10.0.2.15-1610551694034 block blk_107374236
5
```

```
/var/log/hadoop-yarn/apps/root/logs/application_1611089809353_0001/bigdata-srv_4
2232: MISSING 1 blocks of total size 43853 B.Status: CORRUPT
Total size: 627179560 B
```

```
Total dirs:      238
Total files:     260
Total symlinks:      0
Total blocks (validated): 253 (avg. block size 2478970 B)
*****
UNDER MIN REPL'D BLOCKS:      8 (3.1620553 %)
CORRUPT FILES:      8
MISSING BLOCKS:      8
MISSING SIZE:      264814 B
CORRUPT BLOCKS:      8
*****
Corrupt blocks:      8
Number of data-nodes:      1
Number of racks:      1
FSCK ended at Sun Apr 24 21:08:50 PMDT 2022 in 246 milliseconds
```

The filesystem under path '/' is CORRUPT

```
[everis@bigdata-srv script_apoio]$ sudo -u hdfs dfs -rm /tmp/wordcount00002.out/part-r-00000
sudo: dfs: command not found
[everis@bigdata-srv script_apoio]$ everis2021
-bash: everis2021: command not found
[everis@bigdata-srv script_apoio]$ sudo -u hdfs hadoop fsck / | egrep -v '^\.+$' | grep -v eplica
DEPRECATED: Use of this script to execute hdfs command is deprecated.
Instead use the hdfs command for it.
```

```
Connecting to namenode via http://bigdata-srv:50070/fsck?ugi=hdfs&path=%2F
FSCK started by hdfs (auth:SIMPLE) from /192.xxx.x.xxx:x.x for path / at Sun Apr 24 21:14:00
PMDT 2022
/tmp/wordcount00002.out/part-r-00000: CORRUPT blockpool BP-295589643-10.0.2.15-1610551694034
block blk_1073742421
```

```
/tmp/wordcount00002.out/part-r-00000: MISSING 1 blocks of total size 529 B...
/tmp/wordcount012.out/part-r-00000: CORRUPT blockpool BP-295589643-10.0.2.15-1610551694034
block blk_1073742361
```

```
/tmp/wordcount012.out/part-r-00000: MISSING 1 blocks of total size 529 B..
```

```
/user/history/done_intermediate/hdfs/job_1611095949321_0001.summary: CORRUPT blockpool
BP-295589643-10.0.2.15-1610551694034 block blk_1073742422
```

```
/user/history/done_intermediate/hdfs/job_1611095949321_0001.summary: MISSING 1 blocks of
total size 347 B..
/user/history/done_intermediate/root/job_1611089809353_0001-1611091594201-root-word+count-
1611091609627-1-1-SUCCEEDED-root.root-1611091600450.jhist: CORRUPT blockpool
BP-295589643-10.0.2.15-1610551694034 block blk_1073742363
```

```
/user/history/done_intermediate/root/job_1611089809353_0001-1611091594201-root-word+count-
1611091609627-1-1-SUCCEEDED-root.root-1611091600450.jhist: MISSING 1 blocks of total size
33666 B..
/user/history/done_intermediate/root/job_1611089809353_0001.summary: CORRUPT blockpool
BP-295589643-10.0.2.15-1610551694034 block blk_1073742362
```

```
/user/history/done_intermediate/root/job_1611089809353_0001.summary: MISSING 1 blocks of
total size 347 B..
/user/history/done_intermediate/root/job_1611089809353_0001_conf.xml: CORRUPT blockpool
BP-295589643-10.0.2.15-1610551694034 block blk_1073742364
```

/user/history/done_intermediate/root/job_1611089809353_0001_conf.xml: MISSING 1 blocks of total size 141681 B..

/var/log/hadoop-yarn/apps/hdfs/logs/application_1611095949321_0001/bigdata-srv_34788: CORRUPT blockpool BP-295589643-10.0.2.15-1610551694034 block blk_1073742425

/var/log/hadoop-yarn/apps/hdfs/logs/application_1611095949321_0001/bigdata-srv_34788: MISSING 1 blocks of total size 43862 B..

/var/log/hadoop-yarn/apps/root/logs/application_1611089809353_0001/bigdata-srv_42232: CORRUPT blockpool BP-295589643-10.0.2.15-1610551694034 block blk_1073742365

/var/log/hadoop-yarn/apps/root/logs/application_1611089809353_0001/bigdata-srv_42232: MISSING 1 blocks of total size 43853 B.Status: CORRUPT

Total size: 627179560 B

Total dirs: 238

Total files: 260

Total symlinks: 0

Total blocks (validated): 253 (avg. block size 2478970 B)

UNDER MIN REPL'D BLOCKS: 8 (3.1620553 %)

CORRUPT FILES: 8

MISSING BLOCKS: 8

MISSING SIZE: 264814 B

CORRUPT BLOCKS: 8

Corrupt blocks: 8

Number of data-nodes: 1

Number of racks: 1

FSCK ended at Sun Apr 24 21:14:00 PMDT 2022 in 209 milliseconds

The filesystem under path '/' is CORRUPT

[everis@bigdata-srv script_apoio]\$ sudo -u hdfs hdfs dfsadmin -safemode leave

Safe mode is OFF

[everis@bigdata-srv script_apoio]\$ sudo -u hdfs hdfs dfs -rm

/user/history/done_intermediate/root/job_1611089809353_0001-1611091594201-root-word+count-1611091609627-1-1-SUCCEDED-root.root-1611091600450.jhist

Deleted

/user/history/done_intermediate/root/job_1611089809353_0001-1611091594201-root-word+count-1611091609627-1-1-SUCCEDED-root.root-1611091600450.jhist

[everis@bigdata-srv script_apoio]\$ sudo -u hdfs hadoop fsck / | egrep -v '^\.+\$' | grep -v eplica

DEPRECATED: Use of this script to execute hdfs command is deprecated.

Instead use the hdfs command for it.

Connecting to namenode via http://bigdata-srv:50070/fsck?ugi=hdfs&path=%2F

FSCK started by hdfs (auth:SIMPLE) from /192.xxx.x.xxx:x.x for path / at Sun Apr 24 21:23:41 PMDT 2022

/tmp/wordcount00002.out/part-r-00000: CORRUPT blockpool BP-295589643-10.0.2.15-1610551694034 block blk_1073742421

/tmp/wordcount00002.out/part-r-00000: MISSING 1 blocks of total size 529 B...

/tmp/wordcount012.out/part-r-00000: CORRUPT blockpool BP-295589643-10.0.2.15-1610551694034 block blk_1073742361

/tmp/wordcount012.out/part-r-00000: MISSING 1 blocks of total size 529 B..

/user/history/done_intermediate/hdfs/job_1611095949321_0001.summary: CORRUPT blockpool

BP-295589643-10.0.2.15-1610551694034 block blk_1073742422

/user/history/done_intermediate/hdfs/job_1611095949321_0001.summary: MISSING 1 blocks of total size 347 B..

/user/history/done_intermediate/root/job_1611089809353_0001.summary: CORRUPT blockpool BP-295589643-10.0.2.15-1610551694034 block blk_1073742362

/user/history/done_intermediate/root/job_1611089809353_0001.summary: MISSING 1 blocks of total size 347 B..

/user/history/done_intermediate/root/job_1611089809353_0001_conf.xml: CORRUPT blockpool BP-295589643-10.0.2.15-1610551694034 block blk_1073742364

/user/history/done_intermediate/root/job_1611089809353_0001_conf.xml: MISSING 1 blocks of total size 141681 B..

/var/log/hadoop-yarn/apps/hdfs/logs/application_1611095949321_0001/bigdata-srv_34788: CORRUPT blockpool BP-295589643-10.0.2.15-1610551694034 block blk_1073742425

/var/log/hadoop-yarn/apps/hdfs/logs/application_1611095949321_0001/bigdata-srv_34788: MISSING 1 blocks of total size 43862 B..

/var/log/hadoop-yarn/apps/root/logs/application_1611089809353_0001/bigdata-srv_42232: CORRUPT blockpool BP-295589643-10.0.2.15-1610551694034 block blk_1073742365

/var/log/hadoop-yarn/apps/root/logs/application_1611089809353_0001/bigdata-srv_42232: MISSING 1 blocks of total size 43853 B.Status: CORRUPT

Total size: 627146867 B (Total open files size: 380 B)
Total dirs: 237
Total files: 260
Total symlinks: 0 (Files currently being written: 4)
Total blocks (validated): 254 (avg. block size 2469082 B) (Total open file blocks (not validated): 4)

UNDER MIN REPL'D BLOCKS: 7 (2.7559056 %)

CORRUPT FILES: 7

MISSING BLOCKS: 7

MISSING SIZE: 231148 B

CORRUPT BLOCKS: 7

Corrupt blocks: 7

Number of data-nodes: 1

Number of racks: 1

FSCK ended at Sun Apr 24 21:23:41 PMDT 2022 in 55 milliseconds

The filesystem under path '/' is CORRUPT

[everis@bigdata-srv script_apoio]\$ sudo -u hdfs hdfs dfs -rm

/tmp/wordcount012.out/part-r-00000

Deleted /tmp/wordcount012.out/part-r-00000

[everis@bigdata-srv script_apoio]\$ sudo -u hdfs hdfs dfs -rm

/user/history/done_intermediate/hdfs/job_1611095949321_0001.summary

Deleted /user/history/done_intermediate/hdfs/job_1611095949321_0001.summary

[everis@bigdata-srv script_apoio]\$ sudo -u hdfs hdfs dfs -rm

/user/history/done_intermediate/root/job_1611089809353_0001.summary

Deleted /user/history/done_intermediate/root/job_1611089809353_0001.summary

[everis@bigdata-srv script_apoio]\$ sudo -u hdfs hdfs dfs -rm

/user/history/done_intermediate/root/job_1611089809353_0001_conf.xml

Deleted /user/history/done_intermediate/root/job_1611089809353_0001_conf.xml

```
[everis@bigdata-srv script_apoio]$ sudo -u hdfs hdfs dfs -rm
/var/log/hadoop-yarn/apps/hdfs/logs/application_1611095949321_0001/bigdata-srv_34788
Deleted
/var/log/hadoop-yarn/apps/hdfs/logs/application_1611095949321_0001/bigdata-srv_34788
[everis@bigdata-srv script_apoio]$ sudo -u hdfs hdfs dfs -rm
/var/log/hadoop-yarn/apps/root/logs/application_1611089809353_0001/bigdata-srv_42232
Deleted
/var/log/hadoop-yarn/apps/root/logs/application_1611089809353_0001/bigdata-srv_42232
[everis@bigdata-srv script_apoio]$ sudo -u hdfs hadoop fsck / | egrep -v '^\.+$$' | grep -v eplica
DEPRECATED: Use of this script to execute hdfs command is deprecated.
Instead use the hdfs command for it.
```

```
Connecting to namenode via http://bigdata-srv:50070/fsck?ugi=hdfs&path=%2F
FSCK started by hdfs (auth:SIMPLE) from /192.xxx.x.xxx:x.x for path / at Sun Apr 24 21:29:42
PMDT 2022
/tmp/wordcount00002.out/part-r-00000: CORRUPT blockpool BP-295589643-10.0.2.15-1610551694034
block blk_1073742421

/tmp/wordcount00002.out/part-r-00000: MISSING 1 blocks of total size 529 B...
```

```
Status: CORRUPT
Total size:      626915472 B (Total open files size: 166 B)
Total dirs:      237
Total files:      252
Total symlinks:      0 (Files currently being written: 3)
Total blocks (validated):      246 (avg. block size 2548436 B) (Total open file blocks (not
validated): 2)
*****
UNDER MIN REPL'D BLOCKS:      1 (0.40650406 %)
CORRUPT FILES:      1
MISSING BLOCKS:      1
MISSING SIZE:      529 B
CORRUPT BLOCKS:      1
*****
Corrupt blocks:      1
Number of data-nodes:      1
Number of racks:      1
FSCK ended at Sun Apr 24 21:29:42 PMDT 2022 in 85 milliseconds
```

```
The filesystem under path '/' is CORRUPT
[everis@bigdata-srv script_apoio]$ sudo -u hdfs hdfs dfs -rm
/tmp/wordcount00002.out/part-r-00000
Deleted /tmp/wordcount00002.out/part-r-00000
[everis@bigdata-srv script_apoio]$ sudo -u hdfs hadoop fsck / | egrep -v '^\.+$$' | grep -v eplica
DEPRECATED: Use of this script to execute hdfs command is deprecated.
Instead use the hdfs command for it.
```

```
Connecting to namenode via http://bigdata-srv:50070/fsck?ugi=hdfs&path=%2F
FSCK started by hdfs (auth:SIMPLE) from /192.xxx.x.xxx:x.x for path / at Sun Apr 24 21:30:37
PMDT 2022
```

```
Status: HEALTHY
Total size:      626917733 B (Total open files size: 166 B)
Total dirs:      238
```

```
Total files:      254
Total symlinks:    0 (Files currently being written: 3)
Total blocks (validated): 248 (avg. block size 2527894 B) (Total open file blocks (not validated): 2)
Corrupt blocks:    0
Number of data-nodes: 1
Number of racks:   1
FSCK ended at Sun Apr 24 21:30:37 PMDT 2022 in 73 milliseconds
```

```
The filesystem under path '/' is HEALTHY
[everis@bigdata-srv script_apoio]$ hbase shell
2022-04-24 21:35:02,993 INFO [main] Configuration.deprecation: hadoop.native.lib is deprecated. Instead, use io.native.lib.available
HBase Shell; enter 'help<RETURN>' for list of supported commands.
Type "exit<RETURN>" to leave the HBase Shell
Version 1.2.0-cdh5.16.2, rUnknown, Mon Jun 3 03:50:03 PDT 2019
```

```
hbase(main):001:0> status
1 active master, 0 backup masters, 1 servers, 0 dead, 4.0000 average load
```

```
hbase(main):002:0> version
1.2.0-cdh5.16.2, rUnknown, Mon Jun 3 03:50:03 PDT 2019
```

```
hbase(main):003:0> status 'simple'
active master:  bigdata-srv:60000 1650838075363
0 backup masters
1 live servers
    bigdata-srv:60020 1650838074448
        requestsPerSecond=0.0, numberOfOnlineRegions=4, usedHeapMB=26, maxHeapMB=1148, numberOfStores=4, numberOfStorefiles=4, storefileUncompressedSizeMB=0, storefileSizeMB=0, memstoreSizeMB=0, storefileIndexSizeMB=0, readRequestsCount=23, writeRequestsCount=3, rootIndexSizeKB=1, totalStaticIndexSizeKB=0, totalStaticBloomSizeKB=0, totalCompactingKVs=0, currentCompactedKVs=0, compactionProgressPct=NaN, coprocessors=[MultiRowMutationEndpoint]
0 dead servers
Aggregate load: 0, regions: 4
```

```
hbase(main):004:0> status 'summary'
1 active master, 0 backup masters, 1 servers, 0 dead, 4.0000 average load
```

```
hbase(main):005:0> status 'detailed'
version 1.2.0-cdh5.16.2
0 regionsInTransition
active master:  bigdata-srv:60000 1650838075363
0 backup masters
master coprocessors: []
1 live servers
    bigdata-srv:60020 1650838074448
        requestsPerSecond=0.0, numberOfOnlineRegions=4, usedHeapMB=26, maxHeapMB=1148, numberOfStores=4, numberOfStorefiles=4, storefileUncompressedSizeMB=0, storefileSizeMB=0, memstoreSizeMB=0, storefileIndexSizeMB=0, readRequestsCount=23, writeRequestsCount=3, rootIndexSizeKB=1, totalStaticIndexSizeKB=0, totalStaticBloomSizeKB=0, totalCompactingKVs=0, currentCompactedKVs=0, compactionProgressPct=NaN, coprocessors=[MultiRowMutationEndpoint]
        "geog_origin,,1610738538398.849092a95f0f0b2b677f097af4dd3ae7."
        numberOfStores=1, numberOfStorefiles=0, storefileUncompressedSizeMB=
```

```
0, lastMajorCompactionTimestamp=0, storefileSizeMB=0, memstoreSizeMB=0, storefileIndexSizeMB=0, readRequestsCount=0, writeRequestsCount=0, rootIndexSizeKB=0, totalStaticIndexSizeKB=0, totalStaticBloomSizeKB=0, totalCompactingKVs=0, currentCompactedKVs=0, compactionProgressPct=NaN, completeSequenceId=-1, dataLocality=0.0
```

```
"hbase:meta,,1"
```

```
    numberOfStores=1, numberOfStorefiles=2, storefileUncompressedSizeMB=0, lastMajorCompactionTimestamp=1649807933022, storefileSizeMB=0, memstoreSizeMB=0, storefileIndexSizeMB=0, readRequestsCount=19, writeRequestsCount=3, rootIndexSizeKB=1, totalStaticIndexSizeKB=0, totalStaticBloomSizeKB=0, totalCompactingKVs=0, currentCompactedKVs=0, compactionProgressPct=NaN, completeSequenceId=178, dataLocality=1.0
```

```
"hbase:namespace,,1610634073911.28f43b8acea1c80d7a44931f9024464c."
```

```
    numberOfStores=1, numberOfStorefiles=1, storefileUncompressedSizeMB=0, lastMajorCompactionTimestamp=0, storefileSizeMB=0, memstoreSizeMB=0, storefileIndexSizeMB=0, readRequestsCount=4, writeRequestsCount=0, rootIndexSizeKB=0, totalStaticIndexSizeKB=0, totalStaticBloomSizeKB=0, totalCompactingKVs=0, currentCompactedKVs=0, compactionProgressPct=NaN, completeSequenceId=-1, dataLocality=1.0
```

```
"personal,,1610634184168.b04e98ea3a0d648beaba9fe57c701d7f."
```

```
    numberOfStores=1, numberOfStorefiles=1, storefileUncompressedSizeMB=0, lastMajorCompactionTimestamp=0, storefileSizeMB=0, memstoreSizeMB=0, storefileIndexSizeMB=0, readRequestsCount=0, writeRequestsCount=0, rootIndexSizeKB=0, totalStaticIndexSizeKB=0, totalStaticBloomSizeKB=0, totalCompactingKVs=0, currentCompactedKVs=0, compactionProgressPct=NaN, completeSequenceId=-1, dataLocality=1.0
```

```
0 dead servers
```

```
hbase(main):006:0> list  
TABLE  
geog_origin  
personal  
2 row(s) in 0.1450 seconds
```

```
=> ["geog_origin", "personal"]  
hbase(main):007:0> create  
Display all 504 possibilities? (y or n)  
hbase(main):007:0> create
```

```
ERROR: wrong number of arguments (0 for 1)
```

Creates a table. Pass a table name, and a set of column family specifications (at least one), and, optionally, table configuration. Column specification can be a simple string (name), or a dictionary (dictionaries are described below in main help output), necessarily including NAME attribute.
Examples:

```
Create a table with namespace=ns1 and table qualifier=t1  
hbase> create 'ns1:t1', {NAME => 'f1', VERSIONS => 5}
```

```
Create a table with namespace=default and table qualifier=t1  
hbase> create 't1', {NAME => 'f1'}, {NAME => 'f2'}, {NAME => 'f3'}  
hbase> # The above in shorthand would be the following:  
hbase> create 't1', 'f1', 'f2', 'f3'  
hbase> create 't1', {NAME => 'f1', VERSIONS => 1, TTL => 2592000, BLOCKCACHE => true}
```

```

hbase> create 't1', {NAME => 'f1', CONFIGURATION => {'hbase.hstore.blockingStoreFiles' => '10'}}
hbase> create 't1', {NAME => 'f1', IS_MOB => true, MOB_THRESHOLD => 1000000, MOB_COMPACT_PARTITION_POLICY => 'weekly'}

```

Table configuration options can be put at the end.

Examples:

```

hbase> create 'ns1:t1', 'f1', SPLITS => ['10', '20', '30', '40']
hbase> create 't1', 'f1', SPLITS => ['10', '20', '30', '40']
hbase> create 't1', 'f1', SPLITS_FILE => 'splits.txt', OWNER => 'johndoe'
hbase> create 't1', {NAME => 'f1', VERSIONS => 5}, METADATA => { 'mykey' => 'myvalue' }
hbase> # Optionally pre-split the table into NUMREGIONS, using
hbase> # SPLITALGO ("HexStringSplit", "UniformSplit" or classname)
hbase> create 't1', 'f1', {NUMREGIONS => 15, SPLITALGO => 'HexStringSplit'}
hbase> create 't1', 'f1', {NUMREGIONS => 15, SPLITALGO => 'HexStringSplit', REGION_REPLICATION => 2, CONFIGURATION => {'hbase.hregion.scan.loadColumnFamiliesOnDemand' => 'true'}}
hbase> create 't1', {NAME => 'f1', DFS_REPLICATION => 1}

```

You can also keep around a reference to the created table:

```
hbase> t1 = create 't1', 'f1'
```

Which gives you a reference to the table named 't1', on which you can then call methods.

```
hbase(main):008:0> create 'funcionario', 'pessoais', 'profissionais'
0 row(s) in 5.3410 seconds
```

```
=> Hbase::Table - funcionario
```

```
hbase(main):009:0> put 'funcionario', '1', 'pessoais:nome', 'Maria'
0 row(s) in 0.5310 seconds
```

```
hbase(main):010:0> scan 'funcionario'
ROW          COLUMN+CELL
1            column=pessoais:nome, timestamp=1650843647702, value=Maria
1 row(s) in 0.3590 seconds
```

```
hbase(main):011:0> put 'funcionario', '1', 'pessoais:cidade', 'Sao Paulo'
0 row(s) in 0.1290 seconds
```

```
hbase(main):012:0> scan 'funcionario'
ROW          COLUMN+CELL
1            column=pessoais:cidade, timestamp=1650844066348, value=Sao Paulo
1            column=pessoais:nome, timestamp=1650843647702, value=Maria
1 row(s) in 0.0740 seconds
```

```
hbase(main):013:0> put 'funcionario', '2', 'profissionais:empresa', 'Everis'
0 row(s) in 0.1600 seconds
```

```
hbase(main):014:0> scan 'funcionario'
ROW          COLUMN+CELL
1            column=pessoais:cidade, timestamp=1650844066348, value=Sao Paulo
```

```

1          column=pessoais:nome, timestamp=1650843647702, value=Maria
2          column=profissionais:empresa, timestamp=1650844829025,
value=Everis
2 row(s) in 0.0950 seconds

hbase(main):015:0> disable 'funcionario'
0 row(s) in 2.7400 seconds
hbase(main):016:0> alter 'funcionario', NAME=>'hobby', VERSIONS=>5
Updating all regions with the new schema...
1/1 regions updated.
Done.
0 row(s) in 2.8560 seconds

hbase(main):017:0> enable 'funcionario'
0 row(s) in 2.6240 seconds

hbase(main):018:0> scan 'funcionario'
ROW          COLUMN+CELL
1            column=pessoais:cidade, timestamp=1650844066348, value=Sao Paulo
1            column=pessoais:nome, timestamp=1650843647702, value=Maria
2            column=profissionais:empresa, timestamp=1650844829025,
value=Everis
2 row(s) in 0.1650 seconds

hbase(main):019:0> put 'funcionario', '2', 'hobby:nome', 'Ler livros'
0 row(s) in 0.1190 seconds

hbase(main):020:0> scan 'funcionario'
ROW          COLUMN+CELL
1            column=pessoais:cidade, timestamp=1650844066348, value=Sao Paulo
1            column=pessoais:nome, timestamp=1650843647702, value=Maria
2            column=hobby:nome, timestamp=1650845132034, value=Ler livros
2            column=profissionais:empresa, timestamp=1650844829025,
value=Everis
2 row(s) in 0.0340 seconds

hbase(main):021:0> put 'funcionario', '1', 'hobby:nome', 'Ler livros'
0 row(s) in 0.0780 seconds

hbase(main):022:0> scan 'funcionario'
ROW          COLUMN+CELL
1            column=hobby:nome, timestamp=1650845164547, value=Ler livros
1            column=pessoais:cidade, timestamp=1650844066348, value=Sao Paulo
1            column=pessoais:nome, timestamp=1650843647702, value=Maria
2            column=hobby:nome, timestamp=1650845132034, value=Ler livros
2            column=profissionais:empresa, timestamp=1650844829025,
value=Everis
2 row(s) in 0.1930 seconds

hbase(main):023:0> put 'funcionario', '1', 'hobby:nome', 'Pescar'
0 row(s) in 0.1050 seconds

hbase(main):024:0> scan 'funcionario'
ROW          COLUMN+CELL
1            column=hobby:nome, timestamp=1650845195837, value=Pescar
1            column=pessoais:cidade, timestamp=1650844066348, value=Sao Paulo
1            column=pessoais:nome, timestamp=1650843647702, value=Maria

```

```
2          column=hobby:nome, timestamp=1650845132034, value=Ler livros
2          column=profissionais:empresa, timestamp=1650844829025,
value=Everis
2 row(s) in 0.0570 seconds
```

```
hbase(main):025:0> scan 'funcionario', {VERSIONS=>3}
ROW          COLUMN+CELL
1            column=hobby:nome, timestamp=1650845195837, value=Pescar
1            column=hobby:nome, timestamp=1650845164547, value=Ler livros
1            column=pessoais:cidade, timestamp=1650844066348, value=Sao Paulo
1            column=pessoais:nome, timestamp=1650843647702, value=Maria
2            column=hobby:nome, timestamp=1650845132034, value=Ler livros
2            column=profissionais:empresa, timestamp=1650844829025, value=Eve
ris
2 row(s) in 0.1260 seconds
```

```
hbase(main):026:0> count 'funcionario'
2 row(s) in 0.1350 seconds
```

```
=> 2
hbase(main):027:0> scan 'funcionario'
ROW          COLUMN+CELL
1            column=hobby:nome, timestamp=1650845195837, value=Pescar
1            column=pessoais:cidade, timestamp=1650844066348, value=Sao Paulo
1            column=pessoais:nome, timestamp=1650843647702, value=Maria
2            column=hobby:nome, timestamp=1650845132034, value=Ler livros
2            column=profissionais:empresa, timestamp=1650844829025, value=Eve
ris
2 row(s) in 0.0750 seconds
```

```
hbase(main):028:0> delete 'funcionario', '1', 'hobby:nome',
hbase(main):029:0* delete 'funcionario', '1', 'hobby:nome'
SyntaxError: (hbase):29: syntax error, unexpected tSTRING_BEG
```

```
delete 'funcionario', '1', 'hobby:nome'
  ^
```

```
hbase(main):030:0> scan 'funcionario'
ROW          COLUMN+CELL
1            column=hobby:nome, timestamp=1650845195837, value=Pescar
1            column=pessoais:cidade, timestamp=1650844066348, value=Sao Paulo
1            column=pessoais:nome, timestamp=1650843647702, value=Maria
2            column=hobby:nome, timestamp=1650845132034, value=Ler livros
2            column=profissionais:empresa, timestamp=1650844829025, value=Eve
ris
2 row(s) in 0.1550 seconds
```

```
hbase(main):031:0> create 'ttl_exemplo', {'NAME'=>'cf', 'TTL'=>20}
0 row(s) in 2.4080 seconds
```

```
=> Hbase::Table - ttl_exemplo
hbase(main):032:0> put 'ttl_exemplo', 'linha123', 'cf:desc', 'TTL Exemplo'
0 row(s) in 0.1710 seconds
```

```
hbase(main):033:0> scan 'ttl_exemplo'
ROW          COLUMN+CELL
linha123     column=cf:desc, timestamp=1650845745119, value=TTL Exemplo
```

1 row(s) in 0.0630 seconds

hbase(main):034:0> scan 'ttl_exemplo'

ROW COLUMN+CELL

linha123 column=cf:desc, timestamp=1650845745119, value=TTL Exemplo

1 row(s) in 0.0550 seconds

hbase(main):035:0> scan 'ttl_exemplo'

ROW COLUMN+CELL

0 row(s) in 0.0870 seconds

hbase(main):036:0> select * from funcionario

SyntaxError: (hbase):36: syntax error, unexpected tIDENTIFIER

select * from funcionario
^

hbase(main):037:0> Connection reset by 192.xxx.x.xxx:x.x port 22